


Quality vs problem size

median + IQR (shaded only where $M > 4$); 1.0 = brute-force optimum



Runtime vs problem size

median across subsets

